

HW 6 : GANs

In this assignment, you will learn to write an advanced PyTorch implementation concept commonly used in a complex deep learning pipeline by using GAN as a learning example.

You will also start working with more complex architectures (upsampling) and writing style (complex modules such as modulelist) for pytorch.

Highly important -> you should avoid getting K80 GPU in your colab notebook by repeatedly factory resetting the notebook. The K80 is roughly 2x slower than Tesla T4 GPU.

Every TODO is weighted equally. Optional TODO is half of a regular TODO.

GPU test

```
[ ] 1 !nvidia-smi
```

Part 1 : WGAN-GP reimplementation

In this section, you are going to reimplement WGAN-GP (<https://arxiv.org/pdf/1704.00028.pdf>) based on the pseudocode provided in the paper to generate MNIST digit characters. Some parts are intentionally modified to discourage straight copypasting from public repositories.

Algorithm 1 WGAN with gradient penalty. We use default values of $\lambda = 10$, $n_{critic} = 5$, $\alpha = 0.0001$, $\beta_1 = 0$, $\beta_2 = 0.9$.

Require: The gradient penalty coefficient λ , the number of critic iterations per generator iteration n_{critic} , the batch size m , Adam hyperparameters α , β_1 , β_2 .

Require: Initial critic parameters w_0 , initial generator parameters θ_0 .

```
1: while  $\theta$  has not converged do
2:   for  $l = 1, \dots, n_{critic}$  do
3:     for  $i = 1, \dots, m$  do
4:       Sample real data  $x \sim \mathbb{P}_r$ , latent variable  $z \sim p(z)$ , a random number  $\epsilon \sim U[0, 1]$ .
5:        $\tilde{x} \leftarrow G_\theta(z)$ 
6:        $\tilde{x} \leftarrow \epsilon x + (1 - \epsilon)\tilde{x}$ 
7:        $L^{(i)} \leftarrow D_w(\tilde{x}) - D_w(x) + \lambda(\|\nabla_\theta D_w(\tilde{x})\|_2 - 1)^2$ 
8:     end for
9:      $w \leftarrow \text{Adam}(\nabla_w \frac{1}{m} \sum_{i=1}^m L^{(i)}, w, \alpha, \beta_1, \beta_2)$ 
10:   end for
11:   Sample a batch of latent variables  $\{z^{(i)}\}_{i=1}^m \sim p(z)$ .
12:    $\theta \leftarrow \text{Adam}(\nabla_\theta \frac{1}{m} \sum_{i=1}^m -D_w(G_\theta(z)), \theta, \alpha, \beta_1, \beta_2)$ 
13: end while
```

The pseudocode could be organized into two main parts: discriminator optimization in line 2-10, and generator optimization in line 11-12.

The discriminator part consists of four steps:

- Line 4: data, and noise sampling with a batch size of m
- Line 5-7: discriminator loss calculation
- Line 9: discriminator update
- Repeat line 4-9 for n_{critic} steps

After the discriminator is updated, the generator is then updated by performing two steps:

- Line 11: noise sampling
- Line 12: generator loss calculation and update

This part is divided into four subsections: network initialization, hyperparameter initialization, data preparation, and training loop. The detail for each part will be explained in the subsections.

Downloading MNIST dataset

The MNIST dataset contains 60,000 training digit character image (0-9) at 28x28 resolution that are normalized to [0, 1]. Given the training images, your task is to generate new training images using WGAN-GP by learning from the training distribution.

```
1 import torchvision.datasets as datasets
2 import numpy as np
3
4 mnist_trainset = datasets.MNIST(root='./data', train=True, download=True, transform=None)
5 trainX = np.array(mnist_trainset.data[...]).transpose(0, 3, 1, 2) / 255
6 print("Dataset size: ", trainX.shape)
```

Dataset Visualization

```
[ ] 1 import matplotlib.pyplot as plt
2 import numpy as np
3 plt.figure(figsize=(15, 75))
4 for i in range(9):
5     plt.subplot( int('19{}'.format(i+1)) )
6     plt.imshow( trainX[np.random.randint(len(trainX))].transpose((1, 2, 0))[..., 0] , cmap = 'gray' )
7 plt.show()
```

Generator and Discriminator network

Before training, the deep learning networks have to be initialized first. Therefore, in this part, you are going to write a generator and discriminator network based on the description provided below.

The description of the discriminator network is shown in the Table below.

Discriminator (D(x))			
	Kernel size	Resample	Output shape
ConvBlock	5×5	Down	$128 \times 14 \times 14$
ConvBlock	5×5	Down	$256 \times 7 \times 7$
ConvBlock	5×5	Down	$512 \times 4 \times 4$
Linear	-	-	1

The network also has some specific requirements:

- ConvBlock is a Convolution-ReLU layer
- All ReLUs in the encoder are leaky, with a slope of 0.1

The description of the generator network is shown in the Table below.

Generator (G(z))			
	Kernel size	Resample	Output shape
z	-	-	128
Linear	-	-	$512 \times 4 \times 4$
ConvBlock	5×5	Up	$256 \times 8 \times 8$
ConvBlock	5×5	Up	$128 \times 16 \times 16$
Conv, Sigmoid	5×5	Up	$1 \times 32 \times 32$
-	-	Down	$1 \times 28 \times 28$

The network also has some specific requirements:

- ConvBlock is a ConvTranspose-BatchNorm-ReLU layer
- Downsampling method is bilinear interpolation (torch.nn.Upsample or torch.nn.functional.interpolate)

TODO 1: Implement a discriminator network.

TODO 2: Implement a generator network.

```
[ ] 1 import torch
2 import torch.nn.functional as F
3 from torch import nn
4 from torchvision import transforms
5
6 class Discriminator(nn.Module):
7     ##TODO1 implement the discriminator (critic)
8     def __init__(self):
9         super().__init__()
10        self.conv1 = nn.Conv2d(in_channels = 1, out_channels = 128, kernel_size=5, padding=2, stride = 2)
11        self.conv2 = nn.Conv2d(in_channels = 128, out_channels = 256, kernel_size=5, padding=2, stride = 2)
12        self.conv3 = nn.Conv2d(in_channels = 256, out_channels = 512, kernel_size=5, padding=2, stride = 2)
13        self.linear = nn.Linear(512*4*4, 1)
14
15    def forward(self, x):
16        x = self.conv1(x)
17        x = F.leaky_relu(x, negative_slope=0.1)
18        x = self.conv2(x)
19        x = F.leaky_relu(x, negative_slope=0.1)
20        x = self.conv3(x)
21        x = F.leaky_relu(x, negative_slope=0.1)
22        x = x.view(-1, 512*4*4)
23        x = self.linear(x)
24        return x
25
26 class Generator(nn.Module):
27     ##TODO2 implement the generator (actor)
28     def __init__(self):
29         super().__init__()
30        self.linear = nn.Linear(128, 512*4*4)
31        self.conv1 = nn.ConvTranspose2d(in_channels=512, out_channels=256, kernel_size=5, stride=2, padding=2, output_padding=1)
32        self.conv2 = nn.ConvTranspose2d(in_channels=256, out_channels=128, kernel_size=5, stride=2, padding=2, output_padding=1)
33        self.conv3 = nn.ConvTranspose2d(in_channels=128, out_channels=1, kernel_size=5, stride=2, padding=2, output_padding=1)
34        self.upsample = nn.Upsample(size=(28, 28))
35
36    def forward(self, x):
37        # print(x.shape)
38        x = self.linear(x)
39        x = x.view(-1, 512, 4, 4)
40        x = self.conv1(x)
41        x = F.relu(x)
42        x = self.conv2(x)
43        x = F.relu(x)
44        x = self.conv3(x)
45        x = torch.sigmoid(x)
46        x = self.upsample(x)
47        return x
48
49 discriminator = Discriminator().cuda()
50 generator = Generator().cuda()
```

Network verification

TODO 3: What is the input and output shape of the generator and discriminator network? Verify that the implemented networks are the same as the answer you have provided.

```
[ ] 1 !pip install torchinfo
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
```

Parameter Initialization

After the network is initialized, we then set up training hyperparameters for the training. In this part, hyperparameters have already been partially provided in the cell below, though some of them are intentionally left missing (None). Your task is to fill the missing parameters based on the pseudocode above.

TODO4: Initialize the missing model hyperparameters and optimizers based on the pseudocode above.

Note: To hasten the training process of our toy experiment, the training step and batch size is reduced to 3000 and 32, respectively.

```
[ ] 1 NUM_ITERATION = 3000
2 BATCH_SIZE = 32
3 fixed_z = torch.randn((8, 128)).cuda()
4 def schedule(i):
5     lr = 1e-4
6     if(i > 2500): lr *= 0.1
7     return lr
8 losses = {'D': [None], 'G': [None]}
9
10 ## TODO4 initialize missing hyperparameter and optimizer
11 alpha = 0.0001
12 beta1 = 0
13 beta2 = 0.9
14 G_optimizer = torch.optim.Adam(generator.parameters(), lr=alpha, betas=(beta1, beta2))
15 D_optimizer = torch.optim.Adam(discriminator.parameters(), lr=alpha, betas=(beta1, beta2))
16 GP_lambda = 10
17 n_critic = 5
```

• Data preparation

TOD05: Create a dataloader that could generate the data in line 4. The dataloader should return $\mathbf{x}, \mathbf{z}, \mathbf{e}$ with a batch size of `BATCH_SIZE`

```
1 # TOD05: implement dataloader
2 from torch.utils.data import Dataset
3 from torch.utils.data import DataLoader
4 import numpy as np
5 class MinlistDataset(Dataset):
6     def __init__(self, x):
7         self.x = x.astype(float)
8
9     def __getitem__(self, index):
10        x = self.x[index] # Retrieve data
11        # np.random.seed(seed=index)
12        z = np.random.normal(loc=0.0, scale=1.0, size=128)
13        e = np.random.uniform(low=0.0, high=1.0)
14        return x, z, e
15
16    def __len__(self):
17        return self.x.shape[0]
18
19 train_dataset = MinlistDataset(train_x)
20 train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True, pin_memory=True, drop_last=True)
```

• Training loop

This section is the place where the training section starts. It is **highly recommended that you understand the pseudocode before performing the tasks below**. To train the WGAN-GP you have to perform the following tasks:

TOD06: Update the learning rate base on the provided scheduler.

TOD07: Sample the data from the dataloader (Line 4).

TOD08: Calculate the discriminator loss (Line 5-7).

- In the line 7 you have to implement the gradient penalty term $\lambda(|(\nabla_{\mathbf{z}} D_{\theta}(\hat{\mathbf{x}}))|_2 - 1)^2$, which is a custom gradient. You may read <https://pytorch.org/docs/stable/generated/torch.autograd.grad.html> to find how custom gradient is implemented.
- HINT: Gradient norm calculation is still part of the computation graph.

TOD09: Update the discriminator loss (Line 9).

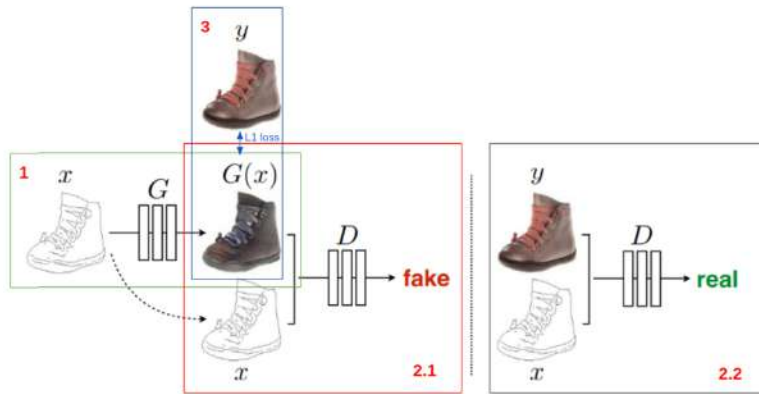
TOD010: Calculate and update the generator loss (Line 11-12).

If your implementation is correct, the generated images should resemble an actual digit character after 500 iterations.

```
1 from tqdm import tqdm
2 model = torch.FloatTensor([1]).cuda()
3 model2 = torch.FloatTensor([-1]).cuda()
4 for i in tqdm(range(NUM_ITERATION)):
5     ## TOD06: update learning rate
6     lr = scheduler(i)
7     for idx in range(len(G_optimizer.param_groups)):
8         G_optimizer.param_groups[idx]['lr'] = lr
9     for idx in range(len(D_optimizer.param_groups)):
10        D_optimizer.param_groups[idx]['lr'] = lr
11
12    for t in range(n_critic):
13        ## TOD07: line 4: sample data from dataloader
14        dataloader_iterator = iter(train_loader)
15        x, z, e = next(dataloader_iterator)
16        ## TOD08: line 5-7: calculate discriminator loss
17        for p in discriminator.parameters():
18            p.requires_grad = True
19        x_hat, e_hat = x.cuda(), z.cuda(), e.cuda()
20        x = x.type(torch.cuda.FloatTensor)
21        x = torch.autograd.Variable(x)
22        discriminator.zero_grad()
23        real_out = discriminator(x).mean()
24
25
26        z = z.type(torch.cuda.FloatTensor)
27        z = torch.autograd.Variable(z, volatile=True)
28        fake = torch.autograd.Variable(generator(z))
29        fake_out = discriminator(fake).mean()
30
31
32        e = e.view(BATCH_SIZE, -1).type(torch.cuda.FloatTensor)
33        e_hat = e[:, None, None].expand(x.shape)
34        x_hat = e*x + (1-e)*fake
35        interpolates = torch.autograd.Variable(x_hat, requires_grad=True)
36        disc_interpolates = discriminator(interpolates)
37        gradients = torch.autograd.grad(outputs=disc_interpolates, inputs=interpolates,
38                                       grad_outputs=torch.ones(disc_interpolates.size()).cuda(),
39                                       create_graph=True, retain_graph=True, only_inputs=True)[0]
40        gradient_penalty = ((gradients.norm(2, dim=1) - 1) ** 2).mean() * GP_lambda
41        loss = fake_out - real_out + gradient_penalty
42        ## TOD09: line 9 update discriminator loss
43        loss.backward()
44        losses['D'].append(loss.mean().item())
45        D_optimizer.step()
46    ## TOD010: line 11-12 calculate and update the generator loss
47
48    for p in discriminator.parameters():
49        p.requires_grad = False
50    generator.zero_grad()
51
52    dataloader_iterator = iter(train_loader)
53    x, z, e = next(dataloader_iterator)
54    x_hat, z_hat, e_hat = x.cuda(), z.cuda(), e.cuda()
55
56    x = x.type(torch.cuda.FloatTensor)
57    x = torch.autograd.Variable(x)
58
59    z = z.type(torch.cuda.FloatTensor)
60    z = torch.autograd.Variable(z)
61
62    fake = generator(z)
63    DoG = -discriminator(fake).mean()
64    DoG.backward()
65    D_optimizer.step()
66
67    losses['G'].append(DoG.item())
68
69
70    # Output visualization: If your reimplementation is correct, the generated images should start resembling a digit character after 500 iterations.
71    if(i % 100 == 0):
72        plt.figure(figsize = (15,75))
73        print(losses['D'][-1], losses['G'][-1])
74        with torch.no_grad():
75            res = generator(fixed_z).cpu().detach().numpy()
76            for k in range(8):
77                plt.subplot( int('18').format(k+1))
78                plt.imshow( res[k].transpose(1, 2, 0)[..., 0], cmap = 'gray' )
79            plt.show()
```

• Part 2 : pix2pix reimplementation (cGAN on paired image translation)

In this exercise, we are reimplementing a paired image translation model, an application of a generative adversarial network (GAN). The model we are going to implement is pix2pix (<https://arxiv.org/pdf/1611.07004.pdf>), one of the earliest paired image translation models based on GAN. The pipeline of pix2pix is shown in the Figure below.



From the figure above, the pipeline consists of three main parts:

- 1. Generation phase: the generator G create the generated image $G(x)$ from the given input x .
- 2. Discrimination phase:
 - In step 2.1, the discriminator D receives an input image x and the generated image $G(x)$, then the discriminator has to learn to predict that the generated image $G(x)$ is fake.
 - In step 2.2, the discriminator D receives an input image x and the ground truth image y , then the discriminator has to learn to predict that the image y is real.
- 3. Refinement phase: Refine the quality of the generated image $G(x)$ by encouraging the generated image to be close to an actual image y by using L1 as an objective.

The objective of pix2pix is to train an optimal generator G^* base on the objective function: $G^* = \arg\min_G \max_D L_{cGAN}(G, D) + \lambda L_1(G)$

- The term $\arg\min_G \max_D L_{cGAN}(G, D)$ is the objective function of the first and second step, which is a standard cGAN loss: $L_{cGAN}(G, D) = E_{x,y}[\log D(x, y)] + E_{x,z}[\log(1 - D(x, G(x, z)))]$. The noise z is embedded in the generator in the form of dropout.
- The term $L_1(G)$ is the objective fuction of the third step where $L_1(G) = E_{x,y,z}[\|y - G(x, z)\|_1]$

The subsections will explain the dataset and training setup of this exercise.

Get dataset

```
[1] 1 wget http://efrosgans.eecs.berkeley.edu/pix2pix/datasets/facades.tar.gz
    2 tar -xzf facades.tar.gz

--2022-04-02 16:20:55-- http://efrosgans.eecs.berkeley.edu/pix2pix/datasets/facades.tar.gz
Resolving efrosgans.eecs.berkeley.edu (efrosgans.eecs.berkeley.edu)... 128.32.244.190
Connecting to efrosgans.eecs.berkeley.edu (efrosgans.eecs.berkeley.edu)|128.32.244.190|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 30168306 (29M) [application/x-gzip]
Saving to: 'facades.tar.gz'

facades.tar.gz 100%[=====] 28.77M 2.71MB/s in 33s

2022-04-02 16:21:28 (896 KB/s) - 'facades.tar.gz' saved [30168306/30168306]
```

Import library

```
[2] 1 import cv2
    2 import glob
    3 import numpy as np
    4 import torch
    5 import torch.nn.functional as F
    6 from torch import nn
    7 from torchvision import transforms
    8 import matplotlib.pyplot as plt
    9 import matplotlib.gridspec as gridspec
```

Setting up facade dataset


The dataset chosen for this exercise is the CMP Facade Database which is a pair of facade images and its segmented component stored in RGB value. The objective of this exercise is to generate a facade given its simplified segmented component. Both input and output is a 256 x 256 RGB image normalized to $[-1, 1]$.

```
[3] 1 train = (np.array([cv2.imread(i) for i in glob.glob('facades/train/*')], dtype = np.float32) / 255).transpose((0, 3, 1, 2))
    2 train = (train - 0.5) * 2 #shift from [0,1] to [-1, 1]
    3 trainX = train[:, :, :, 256:]
    4 trainY = train[:, :, :, :256]
    5
    6 val = (np.array([cv2.imread(i) for i in glob.glob('facades/val/*')], dtype = np.float32) / 255).transpose((0, 3, 1, 2))
    7 val = (val - 0.5) * 2 #shift from [0,1] to [-1, 1]
    8 valX = val[:, :, :, 256:]
    9 valY = val[:, :, :, :256]
    10
    11 print('Input size : {}, Output size = {}'.format(trainX.shape, trainY.shape))

Input size : (400, 3, 256, 256), Output size = (400, 3, 256, 256)
```

Dataset Visualization

```
[4] 1 import matplotlib.pyplot as plt
    2 import numpy as np
    3 plt.figure(figsize = (30,90))
    4 for i in range(4):
    5     idx = np.random.randint(len(trainX))
    6
    7     plt.subplot( int('19{}'.format(2*i+1)) )
    8     plt.title('Input')
    9     plt.imshow( (0.5 + trainX[idx].transpose((1, 2, 0)) + 0.5)[..., ::-1], cmap = 'gray' )
    10    plt.subplot( int('19{}'.format(2*i+2)) )
    11    plt.title('Output')
    12    plt.imshow( (0.5 + trainY[idx].transpose((1, 2, 0)) + 0.5)[..., ::-1], cmap = 'gray' )
    13    plt.show()
```





Note

If you have trouble understanding the instruction provided in this homework or have any ambiguity about the instruction, you could also read the appendix section (section 6.1-6.2) in the paper for a detailed explanation.

Discriminator network

In this section, we are going to implement a discriminator network of $\text{pix}2\text{pix}$. The description of the discriminator network is provided in the Figure below.



The network also has the following specific requirements:

- All convolutions are 4×4 spatial filters
- ConvBlock is a Convolution-InstanceNorm-ReLU layer
- InstanceNorm is not applied to the first C64 layer
- All ReLUs are leaky, with a slope of 0.2

TODO 11: Implement the discriminator network based on the description above.

TODO 12: What should be the size of the input and output of the discriminator for this task? Verify that the input and output of the implemented network are the same as the answer you have provided.

```

23 1 class Discriminator(nn.Module):
24     #TODO011 implement the discriminator network
25     def __init__(self):
26         super().__init__()
27         self.conv1 = nn.Conv2d(in_channels = 6, out_channels = 64, kernel_size=4, padding=1, stride = 2, padding_mode="reflect")
28         self.conv2 = nn.Conv2d(in_channels = 64, out_channels = 128, kernel_size=4, padding=1, stride = 2, padding_mode="reflect")
29         self.conv3 = nn.Conv2d(in_channels = 128, out_channels = 256, kernel_size=4, padding=1, stride = 2, padding_mode="reflect")
30         self.conv4 = nn.Conv2d(in_channels = 256, out_channels = 512, kernel_size=4, padding=1, stride = 1, padding_mode="reflect")
31         self.conv5 = nn.Conv2d(in_channels = 512, out_channels = 1, kernel_size=4, padding=1, stride = 1, padding_mode="reflect")
32
33     def forward(self, x, y):
34         x = torch.cat((x, y), dim=1)
35         x = self.conv1(x)
36         x = F.leaky_relu(x, negative_slope=0.2)
37         x = self.conv2(x)
38         x = F.instance_norm(x)
39         x = F.leaky_relu(x, negative_slope=0.2)
40         x = self.conv3(x)
41         x = F.instance_norm(x)
42         x = F.leaky_relu(x, negative_slope=0.2)
43         x = self.conv4(x)
44         x = F.instance_norm(x)
45         x = F.leaky_relu(x, negative_slope=0.2)
46         x = self.conv5(x)
47         x = torch.sigmoid(x)
48         return x
49
50 discriminator = Discriminator().cuda()
  
```

```

24 1 print(f'The input shape is BATCH_SIZE*6*256*256')
25 2 print(f'The output shape is BATCH_SIZE*1*30*30')
  
```

The input shape is BATCH_SIZE*6*256*256
The output shape is BATCH_SIZE*1*30*30

```

25 1 !pip install torchinfo
26
27 Requirement already satisfied: torchinfo in /usr/local/lib/python3.7/dist-packages (1.6.5)
  
```

```

26 1 #TODO012 verify the discriminator
27 2 from torchinfo import summary
28 3 batch_size = 1
29 4 summary(discriminator, [(batch_size, 3, 256, 256)], [(batch_size, 3, 256, 256)])
  
```

Layer (type:depth-idx)	Output Shape	Param #
Discriminator	1x30x30x1	110,000

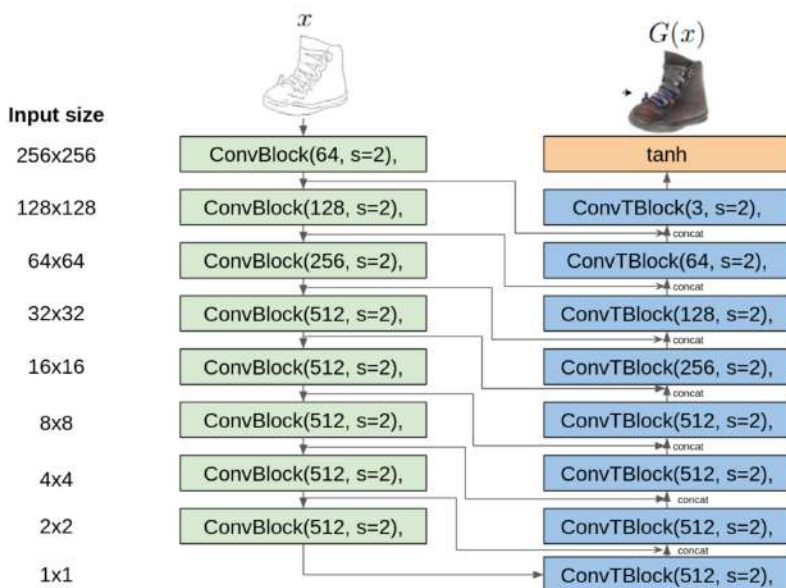
Conv2d: 1-1	[1, 64, 128, 128]	6,208
Conv2d: 1-2	[1, 128, 64, 64]	131,200
Conv2d: 1-3	[1, 256, 32, 32]	524,544
Conv2d: 1-4	[1, 512, 31, 31]	2,097,664
Conv2d: 1-5	[1, 1, 30, 30]	8,193

Total params:	2,767,809	
Trainable params:	2,767,809	
Non-trainable params:	0	
Total mult-adds (G):	3.20	

Input size (MB):	1.57	
Forward/backward pass size (MB):	18.62	
Params size (MB):	11.07	
Estimated Total Size (MB):	31.27	

Generator network

In this section, we are going to implement a generator network of pix2pix. The generator is based on the U-NET based architecture (<https://arxiv.org/abs/1505.04597>). The Description of the generator network is provided in the Figure below.



The network also has the following specific requirements:

- All convolutions are 4 x 4 spatial filters
- ConvBlock is a Convolution-InstanceNorm-ReLU layer
- ConvTBlock is a ConvolutionTranspose-InstanceNorm-Dropout-ReLU layer with a dropout rate of 50%
- InstanceNorm is not applied to the first C64 layer in the encoder
- All ReLUs in the encoder are leaky, with a slope of 0.2, while ReLUs in the decoder are not leaky

TODO 13: Implement the generator network based on the description above.

TODO 14: What should be the size of the input and output of the generator for this task? Verify that the input and output of the implemented network are the same as the answer you have provided.

```

1 #HINT : you could also put multiple layers in a single list using nn.ModuleList
2 class Generator(nn.Module):
3     #TODO13 implement the generator network
4     def __init__(self):
5         super().__init__()
6         in_channels_list = [3, 64, 128, 256, 512, 512, 512, 512]
7         out_channels_list = [64, 128, 256, 512, 512, 512, 512, 512]
8         self.convs = nn.ModuleList([nn.Conv2d(in_channels=in_channels_list[i], out_channels=out_channels_list[i], kernel_size=4, padding=1, stride=2, padding_mode='reflect') for i in range(len(in_channels_list))])
9
10        in_channels_list2 = [512, 1024, 1024, 1024, 1024, 512, 256, 128]
11        out_channels_list2 = [512, 512, 512, 512, 256, 128, 64, 3]
12        self.convTs = nn.ModuleList([nn.ConvTranspose2d(in_channels=in_channels_list2[i], out_channels=out_channels_list2[i], kernel_size=4, stride=2, padding=1) for i in range(len(out_channels_list2))])
13
14    def forward(self, x):
15        send_to_decoder = []
16
17        # Encoder
18        for i in range(len(self.convs)):
19            x = self.convs[i](x)
20            if i > 0 and i < len(self.convs) - 1:
21                x = F.instance_norm(x)
22                x = F.leaky_relu(x, negative_slope=0.2)
23            if i < len(self.convs) - 1:
24                send_to_decoder.append(x)
25
26        send_to_decoder = send_to_decoder[::-1]
27
28        # Decoder
29        for i in range(len(self.convTs)):
30            if i > 0:
31                x = torch.cat([x, send_to_decoder[i-1]], dim=1)
32                x = self.convTs[i](x)
33
34            if i < len(self.convTs) - 1:
35                x = F.instance_norm(x)
36                x = F.dropout(x, p=0.5)
37                x = F.relu(x)
38            x = torch.tanh(x)
39        return x
40
41 generator = Generator().cuda()

```

```

[28] 1 print(f'The input shape is BATCH_SIZE*3*256*256')
2 print(f'The output shape is BATCH_SIZE*3*256*256')

```

The input shape is BATCH_SIZE*3*256*256
The output shape is BATCH_SIZE*3*256*256

```

[29] 1 #TODO14 verify the generator
2 generator.eval()
3 batch_size = 1
4 summary(generator, input_size=(batch_size, 3, 256, 256))

```

Layer (type depth-idx)	Output Shape	Param #
Generator	--	--
ModuleList: 1-1	--	--
ModuleList: 1-2	--	--
ModuleList: 1-1	--	--
Conv2d: 2-1	[1, 64, 128, 128]	3,136
Conv2d: 2-2	[1, 128, 64, 64]	131,200
Conv2d: 2-3	[1, 256, 32, 32]	524,544
Conv2d: 2-4	[1, 512, 16, 16]	2,097,664
Conv2d: 2-5	[1, 512, 8, 8]	4,194,816
Conv2d: 2-6	[1, 512, 4, 4]	4,194,816
Conv2d: 2-7	[1, 512, 2, 2]	4,194,816
Conv2d: 2-8	[1, 512, 1, 1]	4,194,816
ModuleList: 1-2	--	--
ConvTranspose2d: 2-9	[1, 512, 2, 2]	4,194,816
ConvTranspose2d: 2-10	[1, 512, 4, 4]	8,389,120
ConvTranspose2d: 2-11	[1, 512, 8, 8]	8,389,120
ConvTranspose2d: 2-12	[1, 512, 16, 16]	8,389,120
ConvTranspose2d: 2-13	[1, 256, 32, 32]	4,194,560
ConvTranspose2d: 2-14	[1, 128, 64, 64]	1,048,704
ConvTranspose2d: 2-15	[1, 64, 128, 128]	262,208
ConvTranspose2d: 2-16	[1, 3, 256, 256]	6,147
Total params: 54,409,603		
Trainable params: 54,409,603		
Non-trainable params: 0		
Total multi-adds (G): 18.14		
Input size (MB): 0.79		
Forward/backward pass size (MB): 33.72		
Params size (MB): 217.64		
Estimated Total Size (MB): 252.15		

• Data preparation

After the model is initialized, we then create a dataloader to sample the training data. In this paper, to sample the training data, you have to sequentially perform the following steps :

1. Randomly sample the data from the training set
2. Resizing both input and target to 286×286 .
3. Randomly cropping back both images to size 256×256 .
4. Random mirroring the images


TODO15: Implement a dataloader based on the description above. You are allowed to use the function in torchvision.transforms (<https://pytorch.org/vision/main/transforms.html>).


```
[30] 1 # TODO15 implement a dataloader
2 from torch.utils.data import Dataset
3 from torch.utils.data import DataLoader
4 import torchvision.transforms as transforms
5 import numpy as np
6 from PIL import Image
7 import random
8 class FacadeDataset(Dataset):
9     def __init__(self, X, Y):
10         self.x = X.astype(float)
11         self.y = Y.astype(float)
12         self.transform = transforms.Compose([transforms.Resize(size=(286, 286)),\
13                                             transforms.RandomCrop(size=(256, 256)),\
14                                             transforms.RandomHorizontalFlip(p=0.5)])
15
16     def __getitem__(self, index):
17         x = self.x[index] # Retrieve data
18         y = self.y[index]
19         x = torch.from_numpy(x).type(torch.float32)
20         y = torch.from_numpy(y).type(torch.float32)
21         x_and_y = torch.cat([x, y], dim=0)
22         x_and_y = self.transform(x_and_y)
23         x = x_and_y[0, :, :]
24         y = x_and_y[1, :, :]
25         return x, y
26
27     def __len__(self):
28         return self.x.shape[0]
29
30 train_dataset = FacadeDataset(trainX, trainY)
31 val_dataset = FacadeDataset(valX, valY)
32 train_loader = DataLoader(train_dataset, batch_size=1, shuffle=True, pin_memory=True, drop_last=True)
33 val_loader = DataLoader(val_dataset, batch_size=1, shuffle=True, pin_memory=True, drop_last=True)
```

• Dataloader verification

TODO16: Show that the implemented dataloader is working as intended. For instance, are both input and output are flipped and cropped correctly? To obtain a full score, you have to show at least eight data points.

```
[31] 1 # TODO16 show that the dataloader is working properly
2 for i in range(10):
3     dataloader_iterator = iter(train_loader)
4     x, y = next(dataloader_iterator)
5     x, y = x[0], y[0]
6     f, axarr = plt.subplots(1,2)
7     axarr[0].imshow((0.5 * x.detach().numpy().transpose((1, 2, 0)) + 0.5)[..., ::-1], cmap = 'gray')
8     axarr[1].imshow((0.5 * y.detach().numpy().transpose((1, 2, 0)) + 0.5)[..., ::-1], cmap = 'gray')
```





Parameter Initialization

Model hyperparameters and optimizers have already been prepared.

```
[32] 1 import torch.optim as optim
      2 from tqdm import tqdm
      3 lr = 2e-4
      4 LAMBDA = 100
      5 BATCH_SIZE = 1
      6 G_optimizer = optim.Adam(generator.parameters(), lr=lr, betas = (0.5, 0.999))
      7 D_optimizer = optim.Adam(discriminator.parameters(), lr=lr, betas = (0.5, 0.999))
```

Training loop

The training process has the following specific requirements:

- The objective is divided by 2 while optimizing D , which slows down the rate at which D learns relative to G .
- This paper trains the generator G to maximize $\log D(x, G(x, z))$ instead of minimizing $\log(1 - D(x, G(x, z)))$ as the latter term does not provide sufficient gradient.

TODO17: Sample the data using the dataloader.

TODO18: Calculate the discriminator loss and update the discriminator.

- During the update, the loss term $\log(1 - D(x, G(x, z)))$ contains both generator and discriminator. However, we only want to update the discriminator. Therefore, you have to detach the input from the generator graph first. Read <https://pytorch.org/docs/stable/generated/torch.Tensor.detach.html> for additional detail.

TODO19: Calculate the generator loss and update the generator.

HINT

Hint 1: If you are struggling with this part, you could also read the PyTorch DCGAN tutorial as a guideline (https://pytorch.org/tutorials/beginner/dcgan_faces_tutorial.html).

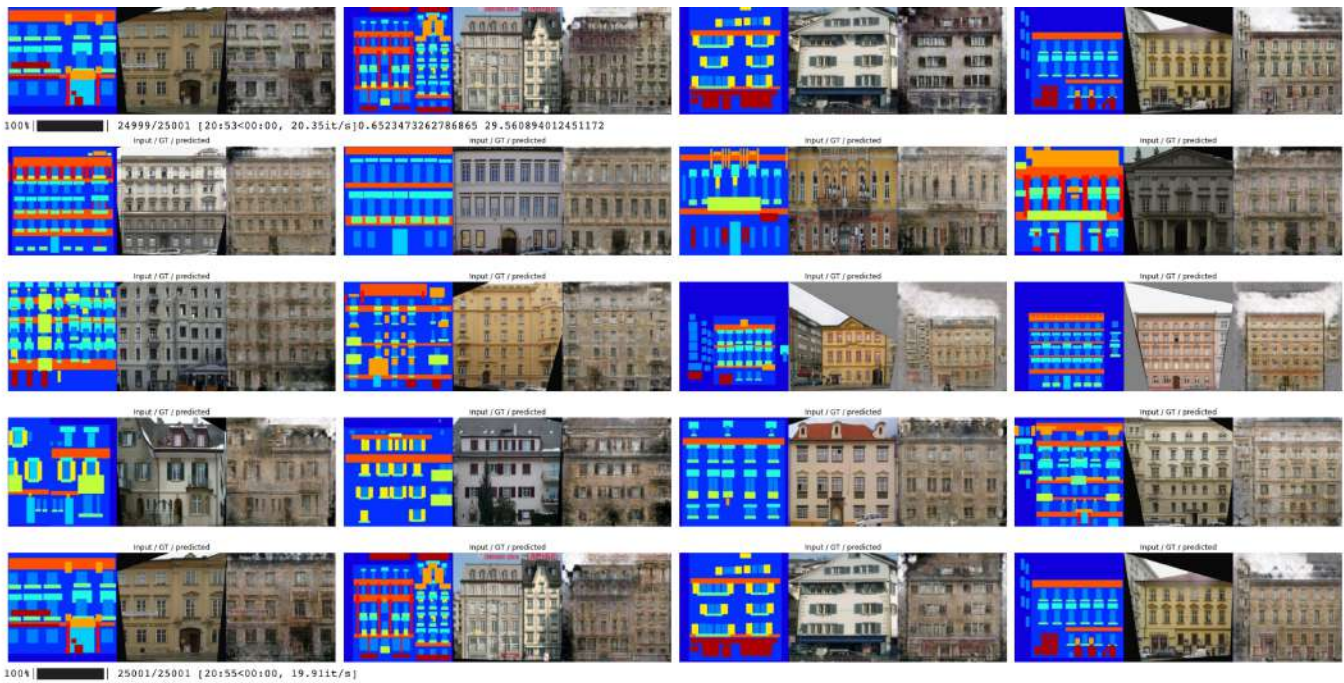
Hint 2: You could remove the L1 loss while debugging since the generator G could still generate the synthetic image even if the L1 loss is removed, though at the cost of increasing image artifacts.

Note


The training schedule in this homework is only one eighth of the original schedule. It is expected that the generated image quality is worse than the one shown in the paper. Nevertheless, the generated facade should still resemble an actual one.

```
[33] 1 losses = {'D': [None], 'G': [None]}
      2 l1_loss = nn.L1Loss()
      3 adversarial_loss = nn.BCELoss()
      4 for i in tqdm(range(25001)):
      5
      6     #TODO17 sample the data from the dataloader
      7     dataloader_iterator = iter(train_loader)
      8     x, y = next(dataloader_iterator)
      9     x, y = x.cuda(), y.cuda()
     10     generator.train()
     11     #TODO18 calculate the discriminator loss and update the discriminator
     12     discriminator.zero_grad()
     13     y_fake = generator(x)
     14     D_real = discriminator(x, y)
     15     D_real_loss = adversarial_loss(D_real, torch.ones_like(D_real))
     16     D_fake = discriminator(x, y_fake.detach())
     17     D_fake_loss = adversarial_loss(D_fake, torch.zeros_like(D_fake))
     18     D_loss = (D_real_loss + D_fake_loss) / 2
     19     losses['D'].append(D_loss.item())
     20     D_loss.backward()
     21     D_optimizer.step()
     22
     23
     24     #TODO19 calculate the generator loss and update the generator
     25     generator.zero_grad()
     26     # generated_img = generator(x)
     27     fake_score = discriminator(x, y_fake)
     28     generator_loss_1 = adversarial_loss(fake_score, torch.ones_like(fake_score).cuda())
     29     generator_loss_2 = LAMBDA * l1_loss(y_fake, y)
     30     # if (i % 100 == 0):
     31     #     print(generator_loss_1, generator_loss_2)
     32     generator_loss = generator_loss_1 + generator_loss_2
     33     losses['G'].append(generator_loss.item())
     34     generator_loss.backward()
     35     G_optimizer.step()
     36     # G_optimizer.update()
     37
     38     # Output visualization : If your reimplementation is correct, the generated images should start resembling a facade after 2,500 iterations
     39     generator.eval()
     40     if (i % 2500 == 0):
     41         with torch.no_grad():
     42             print(losses['D'][-1], losses['G'][-1])
     43             plt.figure(figsize=(40, 16))
     44             gs1 = gridspec.GridSpec(4, 4)
     45             gs1.update(wspace=0.025)
     46
     47             sampleX_vis = 0.5 * valX[:16][:, :, -1, :, :] + 0.5
     48             sampleY = 0.5 * valY[:16][:, :, -1, :, :] + 0.5
     49             sampleX = torch.tensor(valX[:16]).cuda()
     50             pred_val = 0.5 * generator(sampleX.cpu()).detach().numpy()[0][:, :, -1, :, :] + 0.5
     51             vis = np.concatenate([sampleX_vis, sampleY, pred_val], axis=3)
     52             for i in range(vis.shape[0]):
     53                 ax1 = plt.subplot(gs1[i])
     54                 plt.title('Input / GT / predicted')
     55                 plt.axis('off')
     56                 plt.imshow(vis[i].transpose(1, 2, 0))
     57             plt.show()
     58
```






```
[34] 1 plt.plot(np.arange(len(losses['D'])), losses['D'])
      2 plt.title('D loss')
      3 plt.show()
```



```
[35] 1 plt.plot(np.arange(len(losses['G'])), losses['G'])
      2 plt.title('G loss')
      3 plt.show()
```



```
[36] 1 path = "/content/discriminator.pth"
      2 torch.save(discriminator, path)
      3 # /content/facades
      4 path = "/content/generator.pth"
      5 torch.save(generator, path)
```

- (Optional)

Combine the WGAN-GP loss with the pix2pix objective.

```
[36] 1
```